

LECTURE 04D

<EVM>

Nadir Akhtar




BLOCKCHAIN
AT BERKELEY



“WHO AM I?”

Heck if I know

- A blockchain network: →
- Legend:
 - **Nodes:** servers storing blockchain data
 - **Edges:** invisible connection over which peers send and receive data





“99% COMPLETE...”




*low-res aesthetic intentional



“WHO AM I?”

What matters is that you know





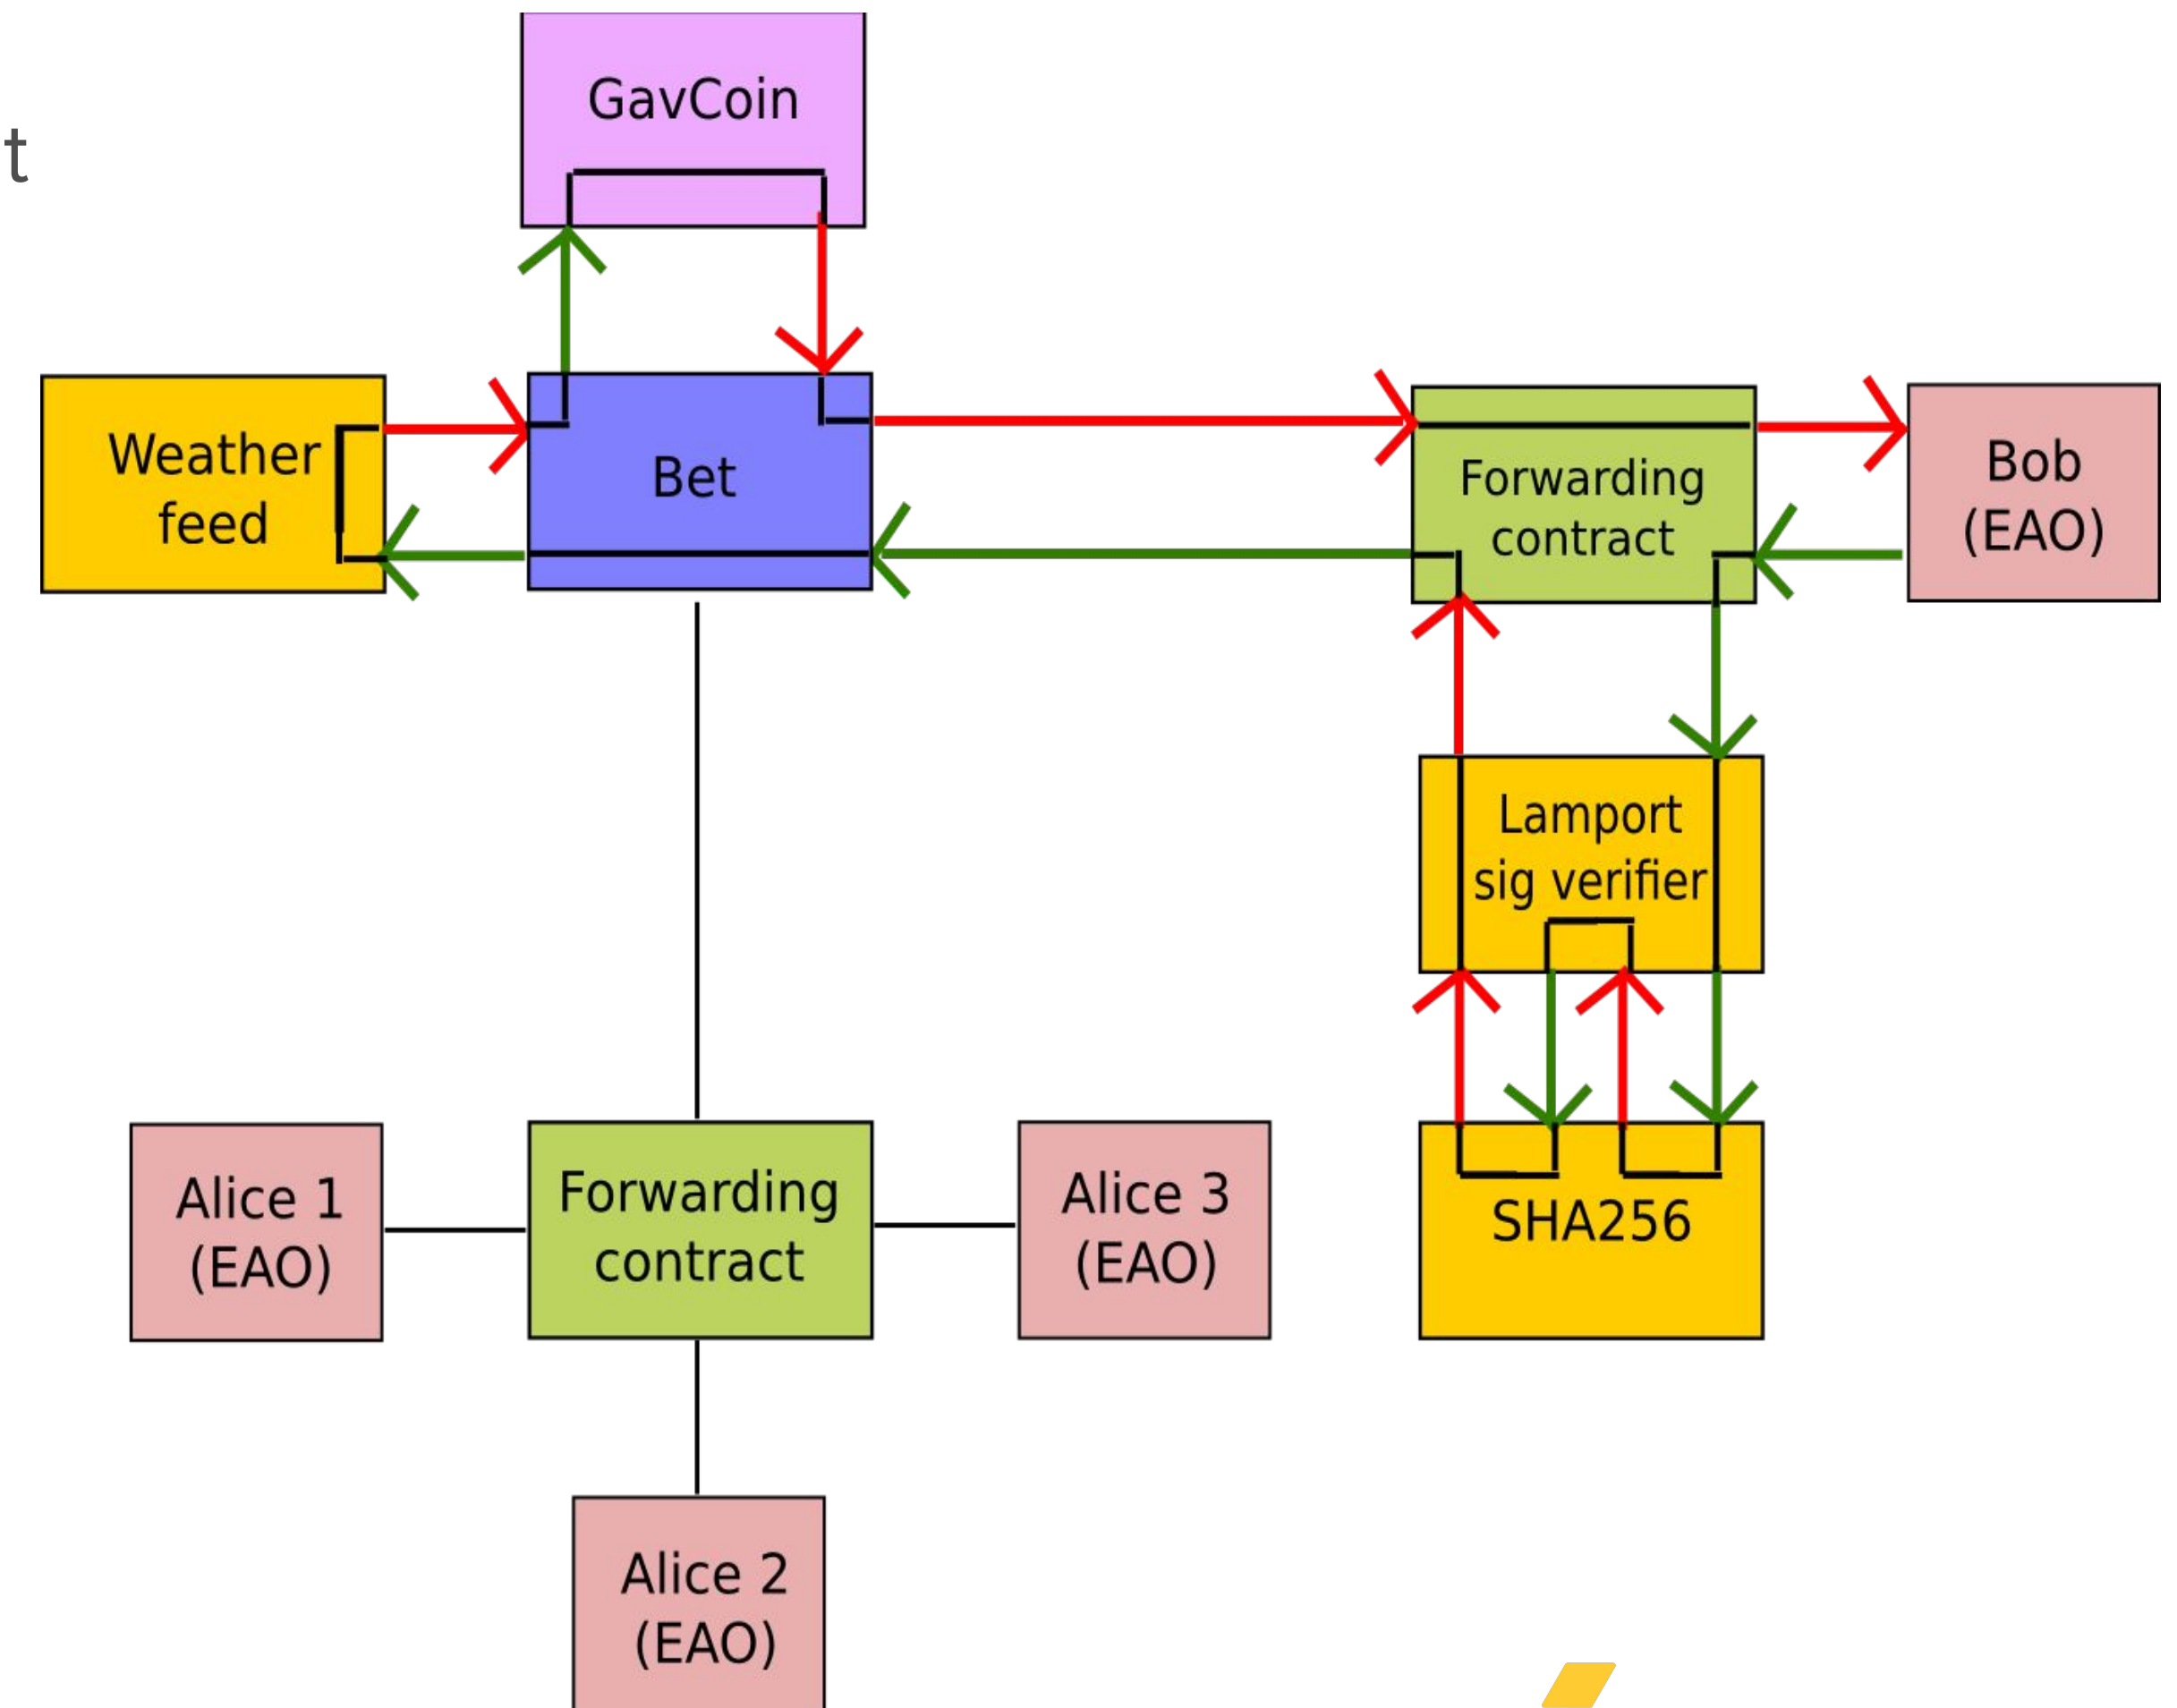
WHAT'S ETHEREUM?

AUTOMATING THE BLOCKCHAIN

<https://github.com/ethereum/wiki/wiki/Ethereum-Development-Tutorial>

- **Smart Contracts**

- Another type of account; not a user account (aka externally owned account (EOA))
- Code with addresses and balances living on the Ethereum blockchain
 - Cannot be edited once deployed
 - Requires no centralization for execution
- Can call each other (and itself, directly or indirectly)
- Run with “gas” to prevent infinite loops
- Use cases: Store and maintain data, create multisig wallet (BitGo), manage contract or user relationships, serve as software library, act as “forwarding contract”





HOW DOES ETHEREUM RUN CODE?


“Is it magic?”

- **EVM (Ethereum Virtual Machine)**


- “[A] large decentralized computer containing millions of objects, called ‘accounts’, which have the ability to maintain an internal database, execute code and talk to each other.”
- Virtual machine: “a software implementation of a machine (i.e. a computer) that executes programs like a physical machine” (think Java VM)
(<http://www.cubrid.org/blog/dev-platform/understanding-jvm-internals/>)
- Runs bytecode with recursive message-sending functionality on every node on the network to verify blocks
- **memory:** byte-array of infinite size
- **program counter:** pointer to current instruction

EVM as a state transition mechanism:


(block_state, gas, memory, transaction, message, code, stack, pc)



where block_state is the global state containing all accounts and includes balances and long-term storage



PUSH1 0 CALLDATALOAD SLOAD NOT PUSH1 9 JUMPI STOP JUMPDEST PUSH1 32 CALLDATALOAD PUSH1 0 CALLDATALOAD SSTORE





EVM'S UPBRINGING

“Do we have to know all this?”

- **EVM Design Goals:**

- ***Simplicity***: op-codes should be as low-level as possible. The number of op-codes should be minimized.
- ***Determinism***: The execution of EVM code should be deterministic; the same input state should always yield the same output state.
- ***Space Efficiency***: EVM assembly should be as compact as possible
- ***Optimization***: Data sizes tend to be bigger; optimize for these larger addresses, read block and transaction data, interact with state, etc
 - easily handle 20-byte addresses (typical public keys) and custom cryptography with 32-byte values, modular arithmetic used in custom cryptography
- ***Security***: it should be easy to come up with a gas cost model for operations that makes the VM non-exploitable



“DO I HAVE TO WRITE CODE WITH OPCODES?”

“That seems dumb.”

- **No! (unless you want to, I suppose)**

- Solidity and Serpent: examples of higher level languages which compile down to same bytecode

- **Solidity**

- Most similar to JavaScript
- Designed specifically for EVM
- Officially supported, unlike Serpent and LLL
- “Statically typed, supports inheritance, libraries, and complex user-defined types”

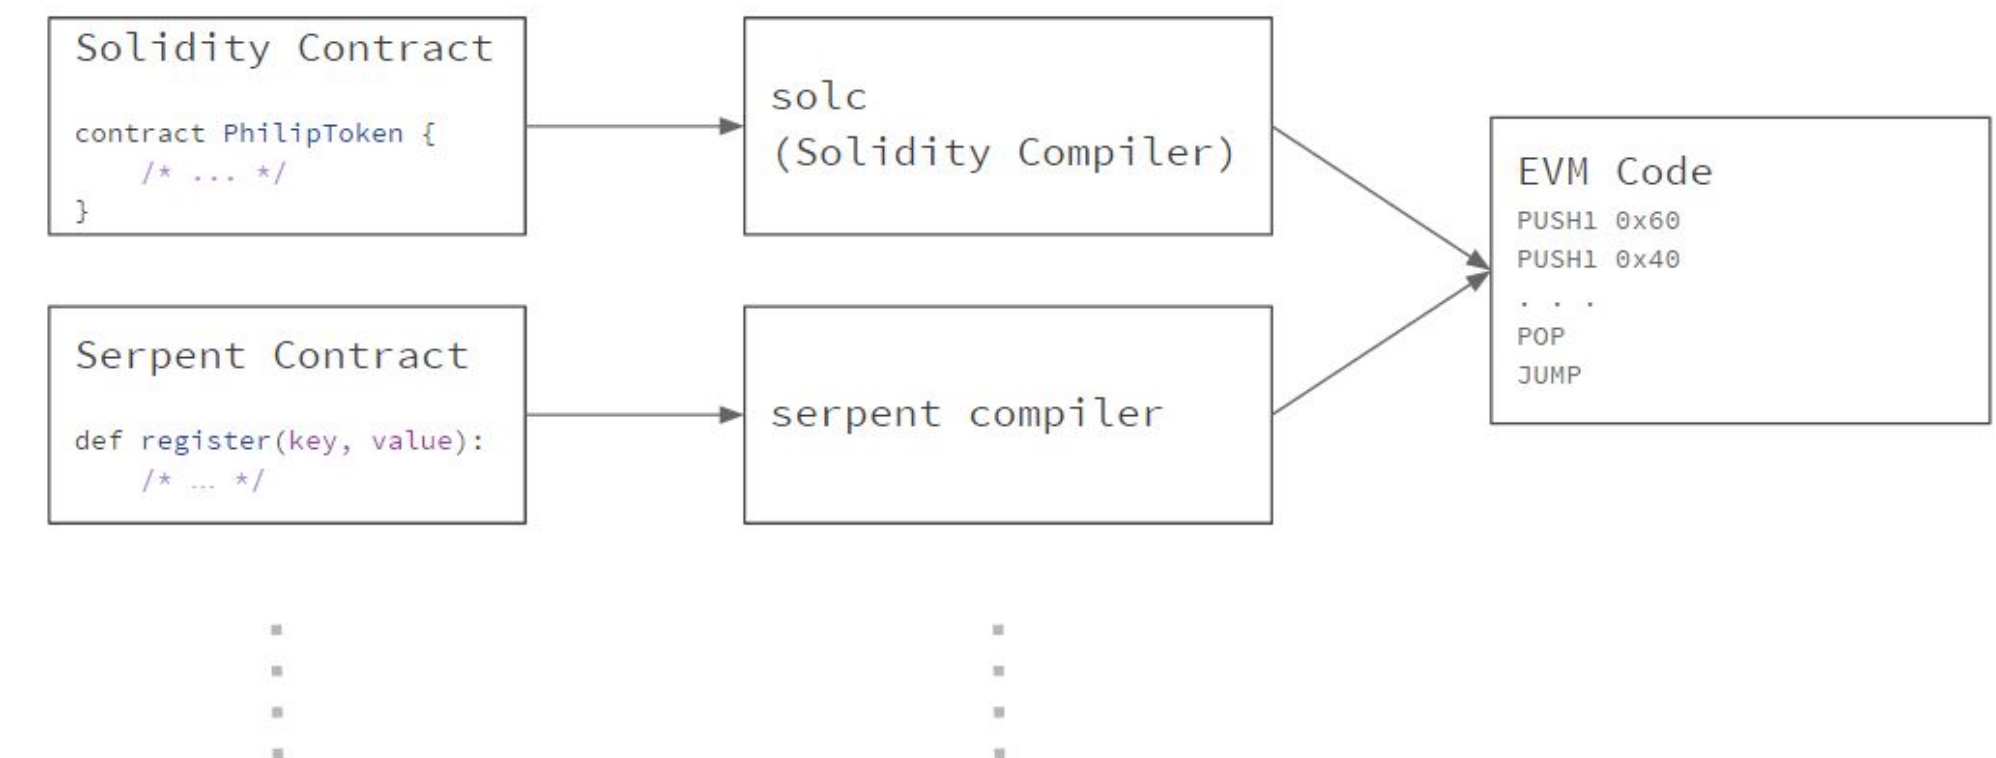
(<https://solidity.readthedocs.io/en/develop/>)

- Most developed language and compiler

- **Compiles into EVM**

- Sent transactions contain EVM bytecode

EVM Code Compilation





EVM GAS

No, nothing to do with bowel movements.

10

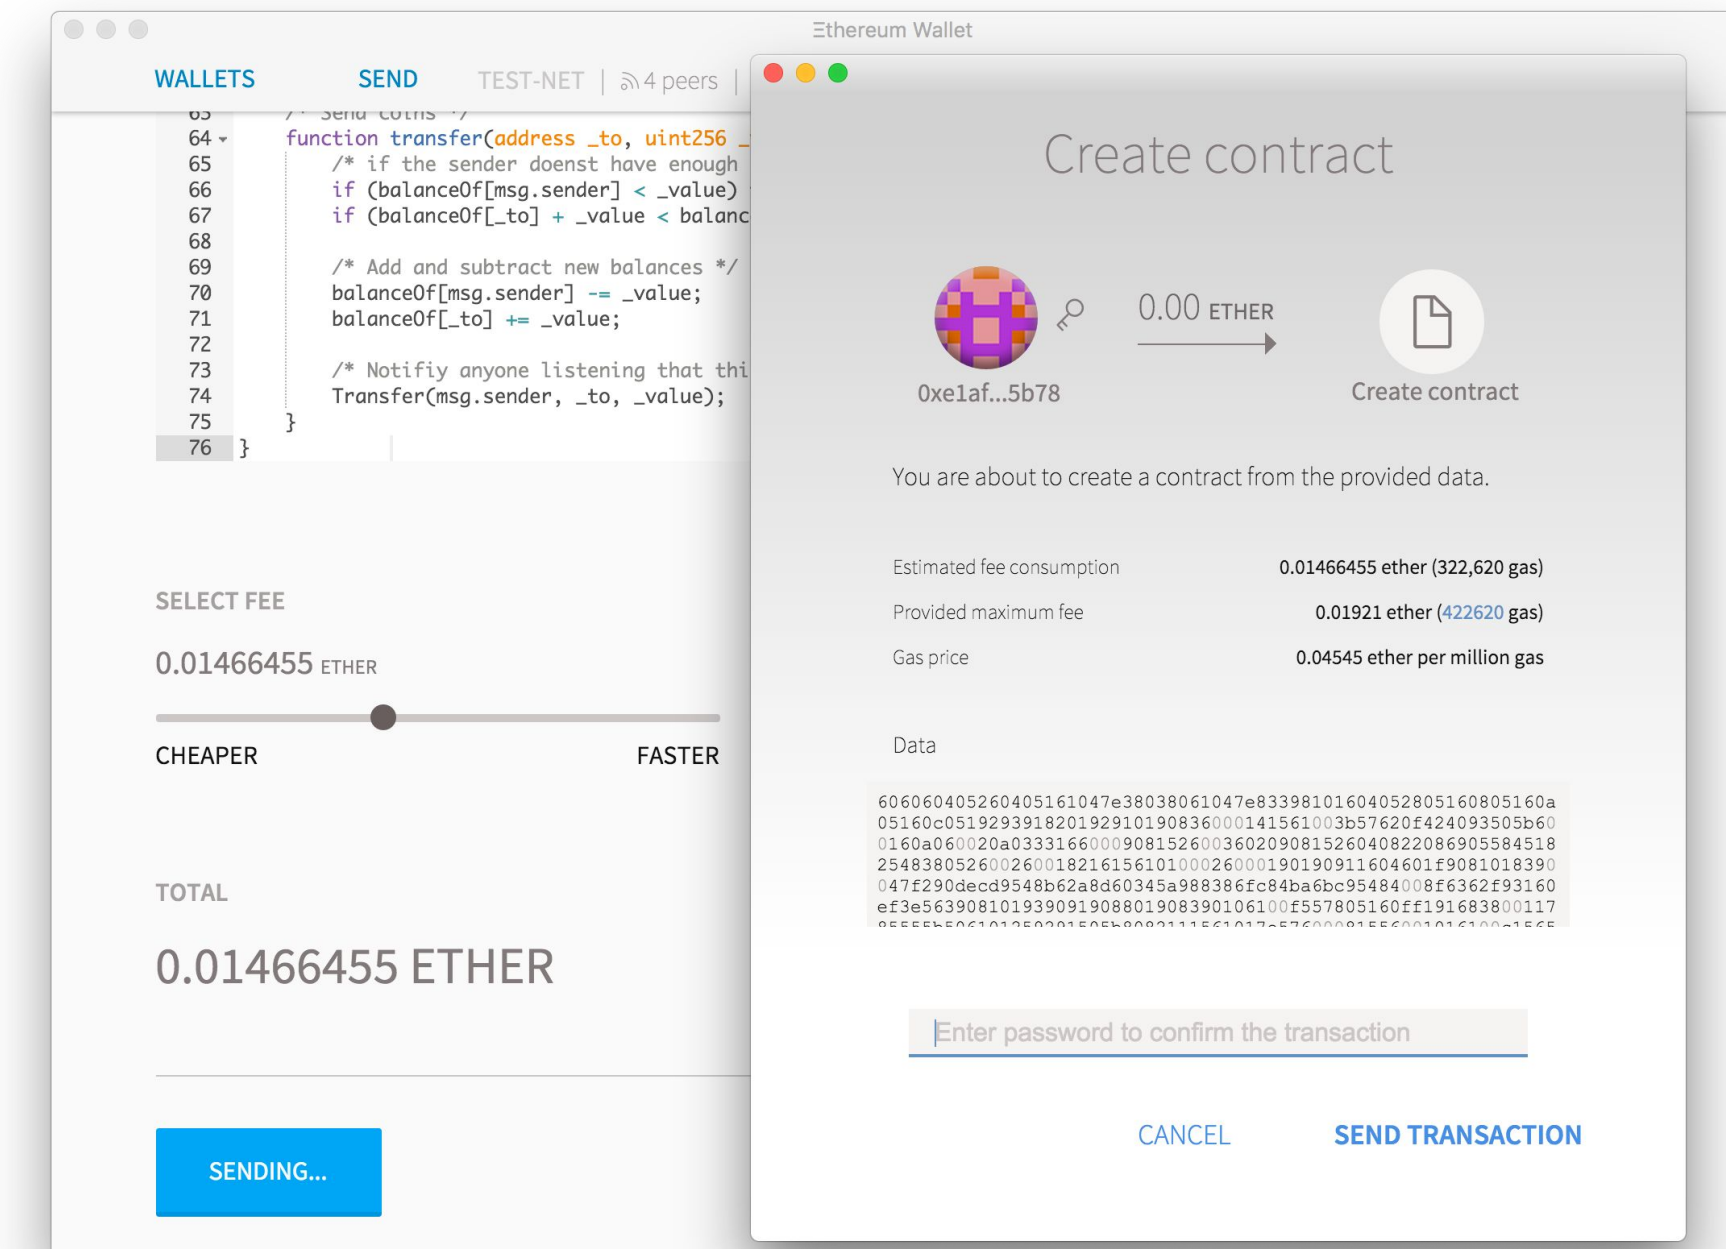
<https://blog.ethereum.org/2015/12/03/how-to-build-your-own-cryptocurrency/>

- **Gas: Ethereum's first defense against attacks**

- Infinite loops are dangerous
- Mathematically impossible to predict infinite loops
- Solution: make contract calls expensive; opcodes eat up money
 - Transaction specifies **startgas** and **gasprice**
 - Miners tend to confirm higher gas prices, similar to Bitcoin TX fees
- Contracts, when calling other contracts, **pass gas** remaining from initial call

- **“Where does the gas go?”**

- Consumed gas goes to miners
- Case a) contract successfully executes:
 - Remaining gas is returned to sender
- Case b) contract runs out of gas before completion or fails to complete
 - Network state reverts, gas is **not refunded**
- Computationally complex code only allowed for high rollers





OPCODE GAS USAGE

EXAMPLES

<https://www.cryptocompare.com/coins/guides/what-is-the-gas-in-et-herium/>

Operation name Gas Cost Function

step	<i>1</i>	Default amount of gas to pay for an execution cycle.
stop	<i>0</i>	Nothing paid for the SUICIDE operation.
sha3	<i>20</i>	Paid for a SHA3 operation.
sload	<i>20</i>	Paid for a SLOAD operation.
sstore	<i>100</i>	Paid for a normal SSTORE operation (doubled or waived sometimes).
balance	<i>20</i>	Paid for a BALANCE operation
create	<i>100</i>	Paid for a CREATE operation
call	<i>20</i>	Paid for a CALL operation.
memory	<i>1</i>	Paid for every additional word when expanding memory
txdata	<i>5</i>	Paid for every byte of data or code for a transaction
transaction	<i>500</i>	Paid for every transaction

PUSH1 0 CALLDATALOAD SLOAD NOT PUSH1 9 JUMPI STOP JUMPDEST PUSH1 32 CALLDATALOAD PUSH1 0 CALLDATALOAD SSTORE



CENTRALIZED VS DECENTRALIZED COMPUTING

IT'S NOT ABOUT THE EFFICIENCY

- **Centralized:**

- Cheaper.
 - No need to do the same operation on thousands of nodes.
- Faster.
 - Information updates much more quickly than any distributed network.
- Single point of failure.
 - Not as open to public regulation or verification, or as robust against attacks or manipulation.

- **Decentralized:**

- Expensive.
 - Smart contract calls no longer of negligible expense
 - Much more time consumed for entire network to update
- Secure.
 - Validated by a whole community.
- Transparent.
 - Anyone can view or participate - including you!

LECTURE 04D

</EVM>

Nadir Akhtar



BLOCKCHAIN
AT BERKELEY